```
ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                      NativeBatchNormBackward0
                                                                    layer1.0.downsample.0.weight (256, 64, 1, 1)
                                                                                                                                       layer1.0.conv1.weight (64, 64, 1, 1)
layer1.0.downsample.1.weight (256) | Layer1.0.downsample.1.bias (256) | AccumulateGrad | MaxPool2DWithIndicesBackward0 | AccumulateGrad | Layer1.0.bn1.weight (64) | Layer1.0.bn1.weigh
                                                                                                                                         ConvolutionBackward0 AccumulateGrad AccumulateGrad
                  AccumulateGrad
                                                AccumulateGrad
                                                                                                                                                  NativeBatchNormBackward0 layer1.0.conv2.weight (64, 64, 3, 3)
                                                        NativeBatchNormBackward0
                                                                                                                                                       ReluBackward0 | layer1.0.bn2.weight (64) | layer1.0.bn2.bias (64)
                                                                                                                                                                           ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                AccumulateGrad
                                                                                                                                                                                     NativeBatchNormBackward0 layer1.0.conv3.weight (256, 64, 1, 1)
                                                                                                                                                                                                                  AccumulateGrad layer1.0.bn3.weight (256) layer1.0.bn3.bias (256)
                                                                                                                                                                                            ReluBackward0
                                                                                                                                                                                                              ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                               AddBackward0 layer1.1.conv1.weight (64, 256, 1, 1)
                                                                                                                                               ReluBackward0 | layer1.1.bn1.weight (64) | layer1.1.bn1.bias (64)
                                                                                                                                                                 ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                            NativeBatchNormBackward0 layer1.1.conv2.weight (64, 64, 3, 3)
                                                                                                                                                                                   ReluBackward0 | layer1.1.bn2.weight (64) | layer1.1.bn2.bias (64)
                                                                                                                                                                                                      ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                               NativeBatchNormBackward0 layer1.1.conv3.weight (256, 64, 1, 1)
                                                                                                                                                                                                                     ReluBackward0 | layer1.1.bn3.weight (256) | layer1.1.bn3.bias (256)
                                                                                                                                                                                                                                        ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                   NativeBatchNormBackward0
                                                                                                                                                                                   AddBackward0 layer1.2.conv1.weight (64, 256, 1, 1)
                                                                                                                                                                                   ReluBackward0 | AccumulateGrad | layer1.2.bn1.weight (64) | layer1.2.bn1.bias (64)
                                                                                                                                                                                                      NativeBatchNormBackward0 layer1.2.conv2.weight (64, 64, 3, 3)
                                                                                                                                                                                                       ReluBackward0 | layer1.2.bn2.weight (64) | layer1.2.bn2.bias (64)
                                                                                                                                                                                                                            ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                       NativeBatchNormBackward0 layer1.2.conv3.weight (256, 64, 1, 1)
                                                                                                                                                                                                                                             ReluBackward0 | AccumulateGrad | layer1.2.bn3.weight (256) | layer1.2.bn3.bias (256)
                                                                                                                                                                                                                                                                ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                            AccumulateGrad ReluBackward0 AccumulateGrad layer2.0.bn1.weight (128) layer2.0.bn1.bias (128)
                                                                                                                                                 AccumulateGrad
                                                                                                                                                                                                                                                      ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                AccumulateGrad
                                                                                                                                                                                         NativeBatchNormBackward0
                                                                                                                                                                                                                                                               NativeBatchNormBackward0 layer2.0.conv2.weight (128, 128, 3, 3)
                                                                                                                                                                                                                                                                      ReluBackward0 | layer2.0.bn2.weight (128) | layer2.0.bn2.bias (128)
                                                                                                                                                                                                                                                                                         ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                 NativeBatchNormBackward0 layer2.0.conv3.weight (512, 128, 1, 1)
                                                                                                                                                                                                                                                                                                        ReluBackward0 | layer2.0.bn3.weight (512) | layer2.0.bn3.bias (512)
                                                                                                                                                                                                                                                                                                                           ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                      NativeBatchNormBackward0
                                                                                                                                                                                                                                                                       AddBackward0 layer2.1.conv1.weight (128, 512, 1, 1)
                                                                                                                                                                                                                                                                      ReluBackward0 | layer2.1.bn1.weight (128) | layer2.1.bn1.bias (128)
                                                                                                                                                                                                                                                                                         ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                 NativeBatchNormBackward0 layer2.1.conv2.weight (128, 128, 3, 3)
                                                                                                                                                                                                                                                                                                         ReluBackward0 | layer2.1.bn2.weight (128) | layer2.1.bn2.bias (128)
                                                                                                                                                                                                                                                                                                                            ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                     NativeBatchNormBackward0 layer2.1.conv3.weight (512, 128, 1, 1)
                                                                                                                                                                                                                                                                                                                                             ReluBackward0 | layer2.1.bn3.weight (512) | layer2.1.bn3.bias (512)
                                                                                                                                                                                                                                                                                                                                                                ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                          NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                        AddBackward0 layer2.2.conv1.weight (128, 512, 1, 1)
                                                                                                                                                                                                                                                                                                          ReluBackward0 | layer2.2.bn1.weight (128) | layer2.2.bn1.bias (128)
                                                                                                                                                                                                                                                                                                                           ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                   AccumulateGrad
                                                                                                                                                                                                                                                                                                                                      NativeBatchNormBackward0 layer2.2.conv2.weight (128, 128, 3, 3)
                                                                                                                                                                                                                                                                                                                                             ReluBackward0 | layer2.2.bn2.weight (128) | layer2.2.bn2.bias (128)
                                                                                                                                                                                                                                                                                                                                                              ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                          NativeBatchNormBackward0 layer2.2.conv3.weight (512, 128, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                ReluBackward0 | layer2.2.bn3.weight (512) | layer2.2.bn3.bias (512)
                                                                                                                                                                                                                                                                                                                                                                                                  ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                             AddBackward0 layer2.3.conv1.weight (128, 512, 1, 1)
                                                                                                                                                                                                                                                                                                                                             ReluBackward0 | AccumulateGrad | layer2.3.bn1.weight (128) | layer2.3.bn1.bias (128)
                                                                                                                                                                                                                                                                                                                                                               ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                              NativeBatchNormBackward0 layer2.3.conv2.weight (128, 128, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                     ReluBackward0 | layer2.3.bn2.weight (128) | layer2.3.bn2.bias (128)
                                                                                                                                                                                                                                                                                                                                                                                      ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                  NativeBatchNormBackward0 layer2.3.conv3.weight (512, 128, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                        ReluBackward0 | layer2.3.bn3.weight (512) | layer2.3.bn3.bias (512)
                                                                                                                                                                                                                                                                                                                                                                                                                        ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                       AccumulateGrad ReluBackward0 AccumulateGrad layer3.0.bn1.weight (256) layer3.0.bn1.bias
                                                                                                                                                                                                                                                                                            layer3.0.downsample.1.weight layer3.0.downsample.1.bias (1024)
                                                                                                                                                                                                                                                                                                                                                                   ConvolutionBackward0
                                                                                                                                                                                                                                                                                                                                           AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                 ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                             AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                         NativeBatchNormBackward0 layer3.0.conv2.weight (256, 256, 3, 3)
                                                                                                                                                                                                                                                                                                                                                   NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                ReluBackward0 | AccumulateGrad | layer3.0.bn2.weight (256) | layer3.0.bn2.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                   ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                             NativeBatchNormBackward0 layer3.0.conv3.weight (1024, 256, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                   ReluBackward0 | layer3.0.bn3.weight (1024) | layer3.0.bn3.bias (1024)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                 ReluBackward0 | AccumulateGrad | layer3.1.bn1.weight (256) | layer3.1.bn1.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                    ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                              NativeBatchNormBackward0 layer3.1.conv2.weight (256, 256, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                   ReluBackward0 | layer3.1.bn2.weight (256) | layer3.1.bn2.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                NativeBatchNormBackward0 layer3.1.conv3.weight (1024, 256, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       ReluBackward0 | AccumulateGrad | layer3.1.bn3.weight (1024) | layer3.1.bn3.bias (1024)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                                                   AddBackward0 layer3.2.conv1.weight (256, 1024, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                     ReluBackward0 | layer3.2.bn1.weight (256) | layer3.2.bn1.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     NativeBatchNormBackward0 layer3.2.conv2.weight (256, 256, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          ReluBackward0 | layer3.2.bn2.weight (256) | layer3.2.bn2.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        NativeBatchNormBackward0 layer3.2.conv3.weight (1024, 256, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             ReluBackward0 | layer3.2.bn3.weight (1024) | layer3.2.bn3.bias (1024)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       AddBackward0 layer3.3.conv1.weight (256, 1024, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        ReluBackward0 | AccumulateGrad | layer3.3.bn1.weight (256) | layer3.3.bn1.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     NativeBatchNormBackward0 layer3.3.conv2.weight (256, 256, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           ReluBackward0 | AccumulateGrad | layer3.3.bn2.weight (256) | layer3.3.bn2.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        NativeBatchNormBackward0 layer3.3.conv3.weight (1024, 256, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              ReluBackward0 | AccumulateGrad | layer3.3.bn3.weight (1024) | layer3.3.bn3.bias (1024)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          AddBackward0 layer3.4.conv1.weight (256, 1024, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            ReluBackward0 | AccumulateGrad | layer3.4.bn1.weight (256) | layer3.4.bn1.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        NativeBatchNormBackward0 layer3.4.conv2.weight (256, 256, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              ReluBackward0 | layer3.4.bn2.weight (256) | layer3.4.bn2.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           NativeBatchNormBackward0 layer3.4.conv3.weight (1024, 256, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                ReluBackward0 | layer3.5.bn1.weight (256) | layer3.5.bn1.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            NativeBatchNormBackward0 layer3.5.conv2.weight (256, 256, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  ReluBackward0 | layer3.5.bn2.weight (256) | layer3.5.bn2.bias (256)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              NativeBatchNormBackward0 layer3.5.conv3.weight (1024, 256, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     ReluBackward0 | layer3.5.bn3.weight (1024) | layer3.5.bn3.bias (1024)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              layer4.0.downsample.1.weight (2048)

layer4.0.downsample.1.bias (2048)

| AccumulateGrad (512) | AccumulateGrad (512) | Compared (512) | Compa
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           ConvolutionBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                NativeBatchNormBackward0 layer4.0.conv2.weight (512, 512, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       ReluBackward0 | layer4.0.bn2.weight (512) | layer4.0.bn2.bias (512)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          ReluBackward0 | AccumulateGrad | layer4.0.bn3.weight (2048) | layer4.0.bn3.bias (2048)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                AddBackward0 layer4.1.conv1.weight (512, 2048, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  ReluBackward0 | AccumulateGrad | layer4.1.bnl.weight | layer4.1.bnl.bias | (512)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   NativeBatchNormBackward0 layer4.1.conv2.weight (512, 512, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         ReluBackward0 | layer4.1.bn2.weight (512) | layer4.1.bn2.bias (512)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      NativeBatchNormBackward0 layer4.1.conv3.weight (2048, 512, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            ReluBackward0 layer4.1.bn3.weight (2048) layer4.1.bn3.bias (2048)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         NativeBatchNormBackward0
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     ReluBackward0 | layer4.2.bn1.weight (512) | layer4.2.bn1.bias (512)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      NativeBatchNormBackward0 layer4.2.conv2.weight (512, 512, 3, 3)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            ReluBackward0 | layer4.2.bn2.weight (512) | layer4.2.bn2.bias (512)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            ConvolutionBackward0 AccumulateGrad AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        NativeBatchNormBackward0 layer4.2.conv3.weight (2048, 512, 1, 1)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              ReluBackward0 | layer4.2.bn3.weight (2048) | layer4.2.bn3.bias (2048)
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  ConvolutionBackward0 AccumulateGrad
                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        AccumulateGrad
```

NativeBatchNormBackward0